

Estimating (non)linear selection on reaction norms
A general framework for labile traits
Introductory tutorial for basic coding and implementation

Jordan S. Martin

Contents

Introduction	2
Generate data	3
Formal model	3
Simulate dataset	4
Code model in Stan	6
Data	6
Parameters and transformed parameters	6
Model	9
Generated quantities	11
Final model code	12
Analyze data	14
Estimate model	14
Investigate results	15
Hypothesis testing	18
Calculate selection gradients	21
Analytic approach (Gaussian case)	21
Numeric approach (general case)	22
References	25

Introduction

This series of tutorials demonstrates how to effectively code and interpret models of nonlinear selection on individual reaction norms (RNs), using the Stan statistical programming language (Carpenter et al. 2017) in R (R Core Team 2020). Stan is an open-source programming language for estimating complex probabilistic models using fully Bayesian inference with state-of-the-art Markov Chain Monte Carlo (MCMC) sampling techniques (Hoffman and Gelman 2014). Stan interfaces with R through the **RStan** and **CmdStan** packages, but you will first need to install Stan on your computer and ensure that it is appropriately configured with your C++ toolchain. Once you are able to effectively use Stan in R, you can begin creating the `.stan` files necessary for estimating the models introduced here. These files can be composed using RStudio or any text editor. A file can also be composed directly in R with `write()`.

```
write("// for Stan comments
      functions{...} // Stan models are composed of
      data {...} // multiple programming blocks
      transformed data {...} //only data, parameters, and model
      parameters {...} //blocks are necessary
      transformed parameters {...}
      model {...}
      generated quantities {...} ",
      "mod1.stan")
```

The `transformed data{}`, `transformed parameters{}`, and `generated quantities{}` blocks are optional and can be used to create additional quantities of interest beyond the initial data provided for the Stan model in `data{}`, the essential model parameters estimated in `parameters{}`, and the likelihood function and priors specified in `model{}`. The utility of these optional blocks will be explored further below. Importantly, any quantities specified in `model{}` will not be saved in the output of the Stan model after estimation. This is very useful for saving memory in a complex analysis.

Once an appropriate `.stan` file is prepared, it can be compiled in R for the C++ toolchain using the `stan_model()` function in RStan or `cmdstan_model()` function in CmdStan and subsequently estimated with an appropriate list of empirical data. The resulting posteriors of a model can then be accessed and manipulated for any further quantities or analyses of interest.

```
#load package
library(rstan) #or
library(cmdstanr)

#make data list
data = list(x = x, y = y, ...)

#compiles the model in C++ for MCMC estimation
mod1 = stan_model("mod1.stan") #or
mod1 = cmdstan_model(stan_file = "mod1.stan", stanc_options = list("O1"))

#samples posterior distribution of the model with default MCMC settings
results = sampling(object = mod1, data = data) #or
results = mod1$sample(data = data)

#extracts posterior estimates
samples = extract(results) #or
samples = results$draws(format = "data.frame")
```

44 Generate data

45 For ease of introduction, this tutorial considers a simple model with a Gaussian phenotype and fitness
 46 measure. As is explained below, the code can be straightforwardly modified for non-Gaussian traits or
 47 fitness components.

48 Formal model

49 It's always helpful to write out the formal model we'd like to estimate in Stan before attempting to code
 50 it. There are a few reasons for this. Firstly, Stan is a probabilistic programming language and, as such,
 51 facilitates coding and estimation of probabilistic models through direct specification of the model parameters
 52 and likelihood functions. Therefore, some understanding of the formal structure of any model is necessary
 53 to flexibly code in Stan. Gaining a deeper understanding of formal statistical models can also be extremely
 54 valuable for building scientists' autonomy and ingenuity in data analysis, which opens up the door to de-
 55 veloping novel models capturing the most salient features of one's specific empirical system and dataset,
 56 rather than being dependent on the assumptions and simplifications of prepackaged toolkits. Researchers
 57 unfamiliar with formal statistical models are encouraged to see McElreath (2020) for a detailed, accessible
 58 treatment in the context of Bayesian data analysis.

59 Notation here follows the main text (see **Eq. 1-2**) for the specific case of normally distributed measurements.
 60 The probability density function for measurement z_{jt} of individual j 's phenotype z at time t is given by

$$\begin{aligned} z_{jt} &\sim \text{Normal}(\mu_{jt}, \sigma_j) \\ \mu_{jt} &= \mu_0 + \mu_{0j} + (\beta_x + \beta_{xj}) x_{jt} \\ \log(\sigma_j^2) &= \sigma_0 + \sigma_{0j} \\ [\mu_0^\top, \beta_x^\top, \sigma_0^\top]^\top &\sim \text{MVNormal}(\mathbf{0}, \mathbf{P}) \end{aligned}$$

61 where $^\top$ indicates the transpose operator. It is important to note here that the (co)variance matrix \mathbf{P} can
 62 be equivalently expressed as the product of a diagonal matrix of standard deviations \mathbf{S} and a correlation
 63 matrix \mathbf{R} , such that $\mathbf{P} = \mathbf{SRS}$. We will exploit this identity below to speed up computation of the model in
 64 Stan. We are also interested in quantifying selection on the individual-level random intercept μ_0 , slope β_x ,
 65 and residual σ_0 parameters. To do so, the full multivariate/multi-response model simultaneously estimates
 66 a probability function for measurement W_{jt} of individual j 's fitness component W at time t

$$\begin{aligned} W_{jt} &\sim \text{Normal}(\theta_{jt}, \delta) \\ \theta_{jt} &= W_0 + W_{0j} + b_1\mu_{0j} + b_2\beta_{xj} + b_3\sigma_{0j} \\ &\quad + q_1\mu_{0j}^2 + q_2\beta_{xj}^2 + q_3\sigma_{0j}^2 + q_4\mu_{0j}\beta_{xj} + q_5\mu_{0j}\sigma_{0j} + q_6\beta_{xj}\sigma_{0j} \\ \mathbf{W}_0 &\sim \text{Normal}(0, \text{sd}(\mathbf{W}_0)) \end{aligned}$$

67 Here \mathbf{b} are regression coefficients quantifying linear effects of the RN parameters on fitness and \mathbf{q} are coeffi-
 68 cients quantifying quadratic effects. Importantly, these coefficients are proportional to but *not* equivalent to
 69 β and γ gradients used in phenotypic selection and quantitative genetic theory. However, we can easily cal-
 70 culate these values by further manipulating the posterior distributions of \mathbf{b} and \mathbf{q} after estimating the model.
 71 This is considered further below. When multiple fitness measures are available for the same subjects, the
 72 model should include individual-level random intercepts W_{0j} to quantify any patterns of repeatable among-
 73 individual differences in fitness that are not accounted for by the modeled phenotype(s) (i.e. unexplained
 74 selection). However, in the absence of repeated fitness measures, variation explained by $\text{sd}(\mathbf{W}_0)$ cannot be
 75 partitioned from variation due to the fitness residual δ and should, therefore, be excluded from the model.

76 Simulate dataset

77 Here we simulate a dataset appropriate for estimating this formal model. To do so, the `NLS_RN_functions.R`
 78 file from the accompanying Github page should be saved locally, providing access to the `sim_RN_Gaus()`
 79 function. By default for this function, population parameter values are fixed so that $\mu_0 = \beta_x = 0$, $\sigma_0 = \log(2)$,
 80 $W_0 = \delta = 1$, and $\text{var}(\mu_0) = \text{var}(\beta_x) = \text{var}(\sigma_0) = \text{var}(W_0) = 1$, with correlations among RN parameters
 81 drawn from an $LKJ(10)$ to produce low to moderate correlations on average. These values can always be
 82 changed by adjusting the corresponding arguments (simply type `sim_RN_Gaus` in the console to reveal the
 83 full function). Using default settings, the user will need to input the range of effect sizes for the selection
 84 effects, which will be drawn from a uniform distribution ranging from `l_es` to `u_es`. For simplicity, these
 85 need to be positive values, but the signs of selection effects will be randomly flipped inside the function. The
 86 user also needs to provide the desired sample size `J`, number of repeated phenotypic measures `rep_z`, and
 87 repeated fitness measures `rep_W`. For the tutorial, we'll generate data for a sample of 1000 individuals with 3
 88 repeated phenotype measures and 2 repeated fitness measures, with selection effects (regression coefficients)
 89 of small-to-moderate statistical effect size.

```
#custom sim function (make sure to set appropriate directory)
#the file should automatically install and load any necessary packages
#that are not found in your R library
source("NLS_RN_functions.R")

#simulate dataset
stan.dl = sim_RN_Gaus(J = 1000, rep_z = 3, rep_W = 2, l_es = 0.2, u_es = 0.3)
```

90 Looking over the structure of this list is useful for understanding how Stan expects data to be formatted.

```
str(stan.dl)
```

```
## List of 11
## $ J      : num 1000
## $ N_z    : num 3000
## $ N_W    : num 2000
## $ ind_z  : int [1:3000] 1 1 1 2 2 2 3 3 3 4 ...
## $ ind_W  : int [1:2000] 1 1 2 2 3 3 4 4 5 5 ...
## $ x      : num [1:3000] 0.332 1.633 -0.801 1.207 -0.357 ...
## $ z      : num [1:3000] -0.114 -3.786 -0.195 -1.555 3.838 ...
## $ W      : num [1:2000] 0.448 -0.214 1.593 3.049 0.306 ...
## $ true_b : num [1:3] 0.251 -0.298 -0.25
## $ true_q : num [1:3] -0.225 -0.277 -0.259
## $ true_qc: num [1:3] 0.217 0.241 -0.208
```

103 The data list not only includes the raw data variables `x` (measured environment), `z` (phenotype), and `W`
 104 (fitness), but also integers `N_z` and `N_W` describing the size of these variables and indices `ind_z` and `ind_W` for
 105 linking repeated observations of phenotype and fitness to individual observations. The integers `N_z` and `N_W`
 106 are used in the Stan model to declare the expected dimensions of `z` and `W`, which ensures that inappropriate
 107 data structures or likelihood functions will throw errors. This will become clearer in the subsequent model
 108 coding section. It is essential that every quantity declared in the `data{}` block of the corresponding Stan
 109 model is also present in this supplied data list. However, the list can also include values of interest that
 110 are not included in the Stan model, such as the true directional `true_b`, stabilizing/disruptive `true_q`, and
 111 correlational selection `true_qc` regression coefficients used for simulating the data. The variables in the data
 112 list that do also appear in `data{}` need to have the same name in both locations.

113 There are a few additional considerations when preparing your own data for using in Stan. There cannot
 114 be any NAs in the supplied list. This means that in some cases, it will be necessary to either drop cases or

115 conduct missing data imputation manually prior to the analysis, or to input an integer that is used within
116 the Stan model to differentiate missing values (e.g. -99 indicates NA). Missing data imputation can also be
117 accomplished during model estimation (see the **Stan Reference Manual** for further details). Character
118 strings can also not be supplied to Stan, including for indexing subject IDs (e.g. “subject1”, “A015”, etc.).
119 Instead, these values will need to be converted to a numeric index. Particular attention needs to be given
120 to the order of data input to the model, as these numeric values will need to be appropriately aligned and
121 indexed throughout the model specification. For instance, if one is changing the character names in an R
122 object `df$subj` for indexing in Stan

```
key.id = unique(df$subj) #all unique subject IDs
new.id = seq(1:length(key.id)) #create numeric index of equal length
df$id = new.id[match(df$subj, key.id)] #numeric id matching order in data frame
#or more simply: df$id = as.integer(factor(df$subj))
```

123 Users will also need to manually ensure that the integers used to index subjects are appropriately aligned
124 with the order of any other data structures corresponding to those subjects. For instance, if one wants to
125 extend the RN model for quantitative genetic analysis using an **A** relatedness matrix, the matrix should be
126 arranged so that row 1 corresponds to the values expected for subject 1 and so on.

```
dimnames(A)[[1]] = new.id[match(dimnames(A)[[1]], key.id)]
dimnames(A)[[2]] = new.id[match(dimnames(A)[[2]], key.id)]
A = as.matrix(A[order(as.numeric(row.names(A))), order(as.numeric(colnames(A)))])
```

127 These additional steps may seem cumbersome at first, but they also allow for many benefits unavailable in
128 more standard statistical software. For example, when each variable can be declared manually with its own
129 separate dimensions and indices, multivariate models can be straightforwardly estimated using data with
130 heterogeneous dimensions (e.g. with differing numbers of subjects and repeated measures for each trait).

Code model in Stan

Stan uses its own language for writing probabilistic models, including a variety of built-in functions designed to aid in efficient computation. The biggest conceptual hurdle for new users of Stan is likely to be the absence of an intuitive R-like syntax for specifying model formulas, such as formulas like $y \sim x + (1|z)$ that can be used to quickly specify complex generalized linear mixed-effects models. These formulas facilitate highly efficient statistical modeling, but do so at the cost of limiting users' ability to specify atypical model structures. Instead, Stan provides the benefit of nearly unlimited flexibility in model specification, with the added cost of a steeper learning curve. In particular, as noted above, models must be formally specified with mathematically appropriate likelihood functions, rather than this process being handled on the back-end through textual inputs from the user such as `family= poisson(link = "log")`. This affords the degree of flexibility necessary for estimating nonlinear selection on RNs. It is important to recognize that some practice and trial-and-error will also be required to gain competency and comfortability with Stan. Therefore, I encourage researchers to review the **Stan Reference Manual**, as well the extensive collection of **Stan Case Studies**, which will provide a more robust foundation for estimating any model of interest in Stan.

As mentioned above, a basic Stan model consists of multiple programming blocks that together specify the data, parameters, likelihood, and quantities of interest for a model. Rather than tackling the model in a single step, we can consider the blocks in turn before putting them together in a single file.

Data

The first component of a Stan model is the data block, which as discussed above tells the model what to expect from our data list, as well as how to treat that data inside the model. Note that Stan uses `//` rather than `#` for comments.

```
data {  
  int<lower=1> J; //total individuals  
  int<lower=1> N_z; //total number of phenotype measures (Z)  
  int<lower=1> N_w; //total number of fitness measures (W)  
  array[N_z] int<lower=1> ind_z; //index of individual measurements (z)  
  array[N_w] int<lower=1> ind_w; //index of individual measurements (W)  
  vector[N_z] x; //environmental covariate  
  vector[N_z] z; //phenotype  
  vector[N_w] W; //fitness  
}
```

The variables in `stan.dl` are declared by their type (note that `int` = integer) with additional arguments regarding their dimensions. `<lower=1>` tells the model to expect sample sizes of at least 1 for the total number of individuals observed `J` and the total number of phenotype `N_z` and fitness measures `N_w`. Our indices `ind_z` and `ind_w`, which link repeated individual observations to the vectors of phenotype `z` and fitness `W` measures, are declared as integers in an array, i.e. multiple integers of length `N_z` and `N_w` that should never be of value less than 1.

Parameters and transformed parameters

The parameters block will take all of the basic parameters that are specified in the nonlinear selection model. Having the formal model above makes this part much easier to code. We can begin by considering the fixed population-level parameters in the model, although the order of specification in the parameters block is entirely arbitrary.

```

parameters {
  //fixed population effects
  real mu_0; //z population intercept
  real beta_x; //z population slope
  real sigma_0; //z population dispersion

  real W_0; //W population intercept
  vector[3] b; //direct selection
  vector[3] q; //stabilizing/disruptive selection
  vector[3] qc; //correlational selection
  //...

```

163 `mu_0` is the population intercept μ_0 for the expectation of behavior `z`, `beta_x` is the population slope β_x ,
 164 and `sigma_0` is the population intercept σ_0 of the dispersion parameter. For the fitness model, we specify
 165 `W_0` for the global intercept W_0 , as well as vectors for the fixed effects describing directional selection `b`
 166 and quadratic selection due to stabilizing/disruptive selection `q` and correlational selection `qc`. In this case,
 167 because there are three RN parameters under selection, there are also 3 corresponding parameters for both
 168 types of quadratic selection. However, in the more general case (e.g. with multiple phenotypes), the vector
 169 `qc` will be of length $\frac{p(p-1)}{2}$ for each unique bivariate combination of p total parameters. Note that we could
 170 also write the stabilizing/disruptive and correlational effects together in a single vector e.g. `vector[3 + 3]`
 171 `q`, or even all selection effects in a single vector e.g. `vector[3+3+3]` `bq`. This choice is entirely arbitrary and
 172 based on what the user finds most intuitive for working with the model.

173 To increase computational efficiency while estimating the random effects, we are going to use a form that is
 174 mathematically equivalent to the formal model above but programmatically distinct. This form combines
 175 two computational tricks—a so-called *non-centered parameterization* of the random effects and a Cholesky
 176 factorization of the random effect correlations—that will help the model to fit much faster. First we'll look
 177 at each step in the code and then consider the maths behind it.

```

//...
//random effects for z
vector<lower=0>[3] sd_RN; //RN parameter sds
matrix[J,3] std_dev_RN; //individual-level RN deviations
cholesky_factor_corr[3] R_chol; //RN parameter correlations

//random effects for W
real<lower=0> sd_W0; //unexplained selection sd
vector[J] std_dev_W; //individual-level selection deviations
real<lower=0> delta; //W dispersion (SD of residuals)
}

```

178 We specify a vector `sd_RN` of length 3 for each of the SDs of the RN parameters and a real number `sd_W0` for
 179 the SD of the individual-level random fitness intercept. The argument `<lower=0>` is essential for letting the
 180 model know that these SDs cannot take on negative values. Given that a variance or SD is necessarily positive
 181 by definition, negative values are mathematically improper solutions and will lead to poor model convergence
 182 as a result. Note that unlike the fitness model, where `delta` is simply the SD of Gaussian fitness residuals,
 183 the dispersion of the phenotype `z` is specified by a linear predictor on the log scale $\log(\sigma_j^2) = \sigma_0 + \sigma_{0j}$, which
 184 is why `sigma_0` is not constrained to be zero (negative values on the log scale are always positive on the
 185 exponentiated scale). A matrix of dimension $(I \times 3)$ `std_dev_RN` is then specified for individuals' deviations
 186 from each of the RN population values (intercept, slope, and residual parameters), and a vector of length I
 187 `std_dev_W` is specified for individuals' deviations from the average fitness intercept. These can be thought
 188 of as standardized best linear unbiased predictors (BLUPs) or simply random effect z-scores.

189 Specifying `std_dev_` and `sd_` as separate parameters allows the model to estimate more efficiently, as the rel-
 190 ative distances between subjects' random effects are separated from (not centered on) the absolute magnitude

of their variation. To see how and why this works, note that any normally distributed random variable

$$\mathbf{v} \sim \text{Normal}(0, \sigma_v)$$

can also be equivalently expressed as a standard normal variable (i.e. a z-score)

$$\mathbf{v}_{\text{std}} \sim \text{Normal}(0, 1)$$

scaled by the original SD

$$\mathbf{v} \equiv \mathbf{v}_{\text{std}} \sigma_v$$

We can use this identity to return the correctly scaled random intercepts \mathbf{W}_0 for the fitness model in the transformed parameters block, which will allow us to use these values directly while specifying the model likelihood.

```
transformed parameters {
  vector[J] W_0j = std_dev_W * sd_W0; //scaled random intercepts for fitness
  //...
```

This code can also be placed in the model block to reduce file size if one isn't interested in saving the scaled values with the model output. To get the scaled random effects for the RN parameters, we also need to bring in the (co)variance between individuals' intercepts, slopes, and residual parameters. This is why we specified a parameter `R_chol` above in the parameters block for the *correlation* matrix \mathbf{R} of the RNs. We do this, rather than specifying a single covariance matrix \mathbf{P} , because we are again going to speed up our computation by separating out the scale of RN parameter deviations (SDs, `sd_RN`) from their standardized associations (correlations), exploiting the identity $\mathbf{P} = \mathbf{SRS}$ explained with the formal model above. In addition, rather than using the Stan function `corr_matrix` for a full correlation matrix, we are also using a special function `Cholesky_factor_corr` to estimate a so-called *Cholesky factorization* of the full \mathbf{R} matrix.

To understand why we do this, note that for any positive definite correlation matrix \mathbf{R} , a Cholesky decomposition can be defined such that

$$\mathbf{R} = \mathbf{R}_L \mathbf{R}_L^\top$$

where \mathbf{R}_L is a lower-triangular matrix and $^\top$ indicates matrix transposition. This property means that we can always estimate the model using a smaller lower-triangular matrix \mathbf{R}_L and subsequently recover the full positive-definite matrix \mathbf{R} by post-multiplying \mathbf{R}_L with its transpose \mathbf{R}_L^\top . This trick is useful for making any Stan model sample more efficiently because it only requires estimating a reduced matrix of lower dimensionality, which lacks the redundant elements of the full symmetric correlation matrix.

With this basic understanding in place, we can now also specify the appropriately scaled random effects for the RN parameters in the transformed parameters block. This is accomplished as follows

```
//...
matrix[J,3] RNj = std_dev_RN * diag_pre_multiply(sd_RN, R_chol)' ;
}
```

where the `*` operator now represents matrix multiplication. The function `diag_pre_multiply` creates a diagonal matrix with `sd_RN`, i.e. $\mathbf{S} = \text{diag}(\text{sd}(\boldsymbol{\mu}_0), \text{sd}(\boldsymbol{\beta}_x), \text{sd}(\boldsymbol{\sigma}_0))$ and then multiplies it with `R_chol`. The product of this multiplication is then transposed $^\top$ using the `'` operator in Stan. This specification gives the appropriate random effects for each individual, as described in the formal model. To see how this works, we can build on the univariate identity above, noting for the multivariate case that with a $(J \times p)$ matrix \mathbf{V} of p phenotypes

$$\mathbf{V} \sim \text{MVNormal}(\mathbf{0}, \mathbf{P})$$

we can derive the appropriately scaled values with a matrix of standard normals \mathbf{V}_{std} and a Cholesky decomposition of \mathbf{P} , so that

$$\mathbf{V} \equiv \mathbf{V}_{\text{std}} \mathbf{P}_L^\top$$

223 where

$$\mathbf{P}_L^\top = \text{Chol}(\mathbf{P})^\top = \text{Chol}(\mathbf{SRS})^\top = (\mathbf{SR}_L)^\top$$

224 In this case, \mathbf{V}_{std} corresponds to `std_dev_RN` and the results of the function `diag_pre_multiply()` represent
225 \mathbf{P}_L as just explained. This may seem like a lot of unnecessary work, but separating out the scale and
226 associations of the random effects in this way will often lead to better model convergence and much more
227 efficient model estimation. Therefore, these mathematically equivalent reparameterizations of the formal
228 model are generally worth implementing although not always strictly necessary. Fortunately, one doesn't
229 need to think much about the maths behind these steps to estimate the model correctly, and this code can
230 generally be copied and applied to any random effect with little modification.

231 Model

232 Now that we've got our data and model parameters specified, it's time to write out the model we'd like to
233 estimate. The model block contains the likelihood functions of the model, the priors for the basic parameters,
234 as well as any data structures that one may want to create for pragmatic convenience in specifying the model
235 but not save in the output (e.g. to reduce memory usage). We'll again work through each component of this
236 block in turn.

```
model{  
  //separate RN parameters  
  vector[J] mu_0j = col(RNj,1); //intercepts  
  vector[J] beta_xj = col(RNj,2); //slopes  
  vector[J] sigma_0j = col(RNj,3); //residuals  
  
  //initialize vectors for response models  
  vector[N_z] mu; //linear predictor of phenotype expectation  
  vector[N_z] sigma; //linear predictor of phenotype dispersion  
  vector[N_W] theta; //linear predictor of fitness expectation  
  //...
```

237 In this first step, we specify a few new vectors to separate out individuals' values for each RN parameter from
238 the matrix `RNj` created in the transformed parameters block. This is a purely pragmatic step that helps to
239 reduce clutter in the model likelihood by avoiding repeated subsetting of the matrix for the respective columns
240 `col(RNj,1)`, `col(RNj,2)`, and `col(RNj,3)`. We do this here rather than in the transformed parameters block
241 because these vectors are redundant with the matrix and would thus be a waste of memory to save in the
242 model output with `RNj`. It is important to realize that the order of indexing the columns in the random
243 effect matrix is arbitrary at this stage. One could also treat column 3 as the intercepts column, for instance,
244 with residuals in column 1. The important thing is that whatever order is used here reflects how the values
245 are specified in the model likelihood. Wherever `mu_0j` and the other parameters are specified in the model
246 likelihood will determine how the values from the respective columns of `RNj` are estimated. To further tidy
247 up the model likelihood code without wasting memory, we create new vectors to temporarily hold the linear
248 predictors of each phenotype and fitness measure. Note that there is no need to create a linear predictor for
249 the dispersion of fitness, as nothing is predicting the residual SD of the fitness model, which is already taken
250 care of by the `delta` parameter.

251 The next step is to fill in these vectors. For the response model of phenotype `z`

```
//...  
//RN model  
mu = mu_0 + mu_0j[ind_z] + (beta_x + beta_xj[ind_z]) .* x;  
sigma = sqrt(exp(sigma_0 + sigma_0j[ind_z]));  
z ~ normal(mu, sigma);
```

252 The index `ind_z` is used to appropriately repeat the random effect values of each RN parameter across
 253 repeated measures of the phenotype. For example, if the first four observations are for individual 1, so that
 254 `ind_z={1,1,1,1,2,...}`, then `mu_0j[ind_z]` will repeat the first value of `mu_0j` for the first four observations
 255 in `z`. This is why it is essential to correctly match the order of the index and the response vectors. The final
 256 line tells Stan that the observed values `z` were generated by a Normal/Gaussian distribution with a likelihood
 257 function described by the expected means `z_mu` and standard deviations `z_sigma` of each observation. Note
 258 that `z_sigma` is calculated with the square root of the exponential function `sqrt(exp())` because the formal
 259 model is specified with a log link function on the variance. We apply the inverse link function to return esti-
 260 mates on the appropriate scale of SDs, i.e. if $\log(\sigma^2) = \sigma_0$ then $\sqrt{\exp(\sigma_0)} = \sigma$. The operator `.*` indicates
 261 element-wise multiplication of vectors, which in this case multiplies the slopes `beta_x + beta_xj[ind_z]`
 262 by the observed environmental measures `x`. These three lines of code are therefore equivalent to

$$\begin{aligned} z_{jt} &\sim \text{Normal}(\mu_{jt}, \sigma_j) \\ \mu_{jt} &= \mu_0 + \mu_{0j} + (\beta_x + \beta_{xj}) x_{jt} \\ \log(\sigma_j^2) &= \sigma_0 + \sigma_{0j} \end{aligned}$$

263 The fitness model can also be specified accordingly

```
//...
//fitness model
theta = W_0 + W_0j[ind_W] +

      b[1] * mu_0j[ind_W] +
      b[2] * beta_xj[ind_W] +
      b[3] * sigma_0j[ind_W] +

      q[1] * (mu_0j[ind_W] .* mu_0j[ind_W]) +
      q[2] * (beta_xj[ind_W] .* beta_xj[ind_W]) +
      q[3] * (sigma_0j[ind_W] .* sigma_0j[ind_W]) +

      qc[1] * (mu_0j[ind_W] .* beta_xj[ind_W]) +
      qc[2] * (mu_0j[ind_W] .* sigma_0j[ind_W]) +
      qc[3] * (beta_xj[ind_W] .* sigma_0j[ind_W]) ;

W ~ normal(theta, delta);
```

264 The `*` operator is used for multiplication with the scalar real values in `b`, `q`, and `qc`, while `.*` is necessary for
 265 element-wise multiplication of the individual values in each vector of RN parameters. This code is equivalent
 266 to our formal fitness model, where we've broken `q` into two separate vectors `q` and `qc` for convenience.

$$\begin{aligned} W_{jt} &\sim \text{Normal}(\theta_{jt}, \delta) \\ \theta_{jt} &= W_0 + W_{0j} + b_1\mu_{0j} + b_2\beta_{xj} + b_3\sigma_{0j} \\ &\quad + q_1\mu_{0j}^2 + q_2\beta_{xj}^2 + q_3\sigma_{0j}^2 + q_4\mu_{0j}\beta_{xj} + q_5\mu_{0j}\sigma_{0j} + q_6\beta_{xj}\sigma_{0j} \end{aligned}$$

267 The final necessary step is to introduce priors for all basic parameters listed in the parameters block. Given
 268 that we're using Bayesian inference for the analysis, specification of the formal model should also include
 269 priors for all estimated model parameters. We encourage using general-purpose, weakly regularizing priors
 270 on model parameters to promote more robust inference and enhance model identification (Lemoine 2019).
 271 The priors we specify for this example will be appropriate for general use with many datasets, but consid-
 272 eration should always be given to the scaling of covariates and intercepts, as well as the desired degree of
 273 regularization. Standardizing variables whenever possible helps to simplify this process.

274 $\mu_0, \beta_x, \sigma_0, W_0, \mathbf{b}, \mathbf{q}, \mathbf{qc} \sim \text{Normal}(0, 1)$
 275 $\text{sd}(\boldsymbol{\mu}_0), \text{sd}(\boldsymbol{\beta}_x), \text{sd}(\boldsymbol{\sigma}_0), \text{sd}(\mathbf{W}_0), \delta \sim \text{Exponential}(2)$
 $\mathbf{R} \sim \text{LKJ}(2)$

276 which translates into Stan code as

```
//...
//model priors

//fixed effects
mu_0 ~ normal(0,1);
beta_x ~ normal(0,1);
sigma_0 ~ normal(0,1);
W_0 ~ normal(0,1);

b ~ normal(0,1);
q ~ normal(0,1);
qc ~ normal(0,1);

//random effects
sd_RN ~ exponential(2);
R_chol ~ lkj_corr_cholesky(2);
to_vector(std_dev_RN) ~ std_normal();

sd_W0 ~ exponential(2);
std_dev_W ~ std_normal();
delta ~ exponential(2);
}
```

277 Note that for the vectors and matrices of random effect deviations `std_dev`, it is necessary to specify that
 278 all elements are described by a `std_normal()` distribution, which makes the non-centered parameterization
 279 introduced in the transformed parameters block above work. The prior `std_normal()` is equivalent to
 280 `normal(0,1)` as written for the other variables, but writing `std_normal()` is helpful to distinguish these
 281 essential priors from those that can be modified based on the goals of the analysis.

282 Generated quantities

283 The final programming block in our Stan model concerns the calculation of any quantities of interest which
 284 weren't directly estimated in earlier blocks. Here we can overcome the complexity of the reparameterizations
 285 used for increased efficiency by saving much easier to interpret full (co)variance \mathbf{P} and correlation matrices
 286 \mathbf{R} , as well as the variances of the RN parameters.

```
generated quantities{
  matrix[3,3] R = R_chol * R_chol'; //RN correlation matrix
  matrix[3,3] S = diag_matrix(sd_RN); //RN SD matrix
  matrix[3,3] P = S*R*S; //RN covariance matrix
  vector<lower=0>[3] V_P = sd_RN .* sd_RN; //RN variances
}
```

287 Final model code

288 With each programming block coded, we can put them all together and write to a single `.stan` file in R.
 289 Doing this in RStudio using a formatted Stan file (File > New File > Stan file) will make coding and
 290 debugging (click Check on save) much easier.

```

write("
data {
  int<lower=1> J; //total individuals
  int<lower=1> N_z; //total number of pheontype meausres (Z)
  int<lower=1> N_W; //total number of fitness measures (W)
  array[N_z] int<lower=1> ind_z; //index of individual measurements (z)
  array[N_W] int<lower=1> ind_W; //index of individual measurements (W)
  vector[N_z] x; //environmental covariate
  vector[N_z] z; //phenotype
  vector[N_W] W; //fitness
}

parameters {
  //fixed population effects
  real mu_0; //z population intercept
  real beta_x; //z population slope
  real sigma_0; //z population dispersion

  real W_0; //W population intercept
  vector[3] b; //direct selection
  vector[3] q; //stabilizing/disruptive selection
  vector[3] qc; //correlational selection

  //random effects for z
  vector<lower=0>[3] sd_RN; //RN parameter sds
  matrix[J,3] std_dev_RN; //individual-level RN deviations
  cholesky_factor_corr[3] R_chol; //RN parameter correlations

  //random effects for W
  real<lower=0> sd_W0; //unexplained selection sd
  vector[J] std_dev_W; //individual-level selection deviations
  real<lower=0> delta; //W dispersion (SD of residuals)
}

transformed parameters {
  vector[J] W_0j = std_dev_W * sd_W0; //scaled random intercepts for fitness
  matrix[J,3] RNj = std_dev_RN * diag_pre_multiply(sd_RN, R_chol)' ;
}

model{
  //separate RN parameters
  vector[J] mu_0j = col(RNj,1); //intercepts
  vector[J] beta_xj = col(RNj,2); //slopes
  vector[J] sigma_0j = col(RNj,3); //residuals

  //initialize vectors for response models
  vector[N_z] mu; //linear predictor of phenotype expectation
  vector[N_z] sigma; //linear predictor of phenotype dispersion

```

```

vector[N_W] theta; //linear predictor of fitness expectation

//RN model
mu = mu_0 + mu_0j[ind_z] + (beta_x + beta_xj[ind_z]) .* x;
sigma = sqrt(exp(sigma_0 + sigma_0j[ind_z]));
z ~ normal(mu, sigma);

//fitness model
theta = W_0 + W_0j[ind_W] +

      b[1] * mu_0j[ind_W] +
      b[2] * beta_xj[ind_W] +
      b[3] * sigma_0j[ind_W] +

      q[1] * (mu_0j[ind_W] .* mu_0j[ind_W]) +
      q[2] * (beta_xj[ind_W] .* beta_xj[ind_W]) +
      q[3] * (sigma_0j[ind_W] .* sigma_0j[ind_W]) +

      qc[1] * (mu_0j[ind_W] .* beta_xj[ind_W]) +
      qc[2] * (mu_0j[ind_W] .* sigma_0j[ind_W]) +
      qc[3] * (beta_xj[ind_W] .* sigma_0j[ind_W]) ;
W ~ normal(theta, delta);

//model priors

//fixed effects
mu_0 ~ normal(0,1);
beta_x ~ normal(0,1);
sigma_0 ~ normal(0,1);
W_0 ~ normal(0,1);

b ~ normal(0,1);
q ~ normal(0,1);
qc ~ normal(0,1);

//random effects
sd_RN ~ exponential(2);
R_chol ~ lkj_corr_cholesky(2);
to_vector(std_dev_RN) ~ std_normal();

sd_W0 ~ exponential(2);
std_dev_W ~ std_normal();
delta ~ exponential(2);
}

generated quantities{
  matrix[3,3] R = R_chol * R_chol'; //RN correlation matrix
  matrix[3,3] S = diag_matrix(sd_RN); //RN SD matrix
  matrix[3,3] P = S*R*S; //RN covariance matrix
  vector<lower=0>[3] V_P = sd_RN .* sd_RN; //RN variances
}

", "mod1.stan")

```

Analyze data

Estimate model

To estimate this model, we first pass it to Stan for C++ compilation. This step will be the same regardless of how the model was coded.

```
#load package
library(rstan)

#rstan: compiles the model in C++ for MCMC estimation
mod1 = stan_model("mod1.stan")

#basic settings for rstan
options(mc.cores = parallel::detectCores())
rstan_options(auto_write = TRUE)

#or in cmdstan
set_cmdstan_path(...) #... = directory of installation
mod1 = cmdstan_model(stan_file = "mod1.stan")
```

The compiled model in `mod1` is now ready to be sampled immediately using Stan's cutting-edge Markov Chain Monte Carlo (MCMC) algorithm, which is accomplished by passing it to the `sampling()` function from RStan. As noted above, the CmdStan package can also be used for more efficient computation. We'll use RStan here because it is less prone to complications during installation and provides more user-friendly functionality. For default MCMC settings in Stan, we could run

```
#sampling posterior dist of the model with default MCMC settings
results = sampling(object = mod1, data = stan.d1)

#or in cmdstan
results = mod1$sample(data = fishd1)
```

However, given that our model is somewhat complex, it is helpful to use custom settings for the sampler that will reduce the risk of poor performance. In particular, we can manually specify that the MCMC sampler should use 500 iterations per chain to converge on the target joint posterior distribution `warmup=500`, with the subsequent 2000 iterations/chain used as posterior samples `iter = 2500` (i.e. `iter - warmup` = number of MCMC samples per chain). For users familiar with Gibbs and Metropolis-Hastings MCMC algorithms, this will likely seem like an insufficient number of samples. However, the No U-Turn sampler implemented in Stan tends to be dramatically more efficient than these classical algorithms, requiring much fewer samples to reach effective sample sizes for accurately approximating posterior distributions. The `init = 0` argument can be used to initialize the samplers near null values, which is not necessary but can aid the speed of convergence for complex models. We'll use four MCMC chains to assess model convergence across independent random samplers `chains=4`, with one core assigned to each chain for parallel processing `cores=4`. The appropriate number of cores to use will be contingent on one's hardware. The `adapt_delta=0.80` argument reduces the risk of divergent transitions during sampling. If errors are thrown about divergence transitions, this value can be increased up to a maximum of 1, e.g. 0.90 or 0.95.

```
#progress of MCMC chains can be tracked in the viewer pane of RStudio
results = sampling(object = mod1, data = stan.d1, warmup=500, iter = 1000, init = 1e-4,
                  chains=4, cores=4, control=list(adapt_delta=0.90) )
```

```
#or in cmdstan (generally faster)
results = mod1$sample(data = stan.dl, iter_warmup = 500, iter_sampling = 1000, init = 1e-4,
                      chains = 4, parallel_chains = 4, adapt_delta = 0.90)
```

Some readers may note that there is no argument specified for thinning the chain, which implicitly specifies the default argument `thin=1`. Although there are specific contexts where thinning is useful for MCMC sampling, it is generally unnecessary and computationally inefficient (Link and Eaton 2012).

If you estimate a model in Stan and receive a warning or error, it may indicate issues with the MCMC sampler, which should always be taken seriously. Further description of these and other warnings can be found in the **Stan Warning Guide**. Some warnings can be safely ignored in particular contexts, but efforts should always be taken to first remove the issue before interpreting or reporting results from the sampler. If you receive a warning regarding divergent transitions, a straightforward first step is to increase the `adapt_delta` value closer to 1, e.g. from 0.95 to 0.99. The higher this value, the slower the model will sample but the less likely that divergent iterations will occur. Similarly, if warnings of bulk or tail ESS are received, a first step is to simply let the chains sample for longer by increasing the `iter`, e.g. from 3000 to 3500 or 4000.

Investigate results

Assuming that the sampling procedure worked as intended, we can move ahead with extracting the posterior MCMC samples from the model.

```
#extracts posterior estimates
samples = rstan::extract(results)

#or cmdstan (using custom function to mimic rstan)
samples = extract(results)


#quick glance of MCMC samples for linear selection effects
head(samples$b[,1:3])
```

```
##
## iterations      [,1]      [,2]      [,3]
##      [1,] -0.2214934 -0.252121794 0.3542963
##      [2,] -0.1978726 -0.111582286 0.3732336
##      [3,] -0.1841798 -0.001276365 0.3766329
##      [4,] -0.1821944 -0.090781450 0.5194939
##      [5,] -0.1951355 -0.064803276 0.3471871
##      [6,] -0.1893245 -0.185388096 0.2622879
```

Before hypothesis testing, it is useful to visualize the shapes and locations of the posterior distributions of model parameters. There are many ways this can be accomplished. For example, the `bayesplot` package can be used to generate a variety of useful plots.

```
library(bayesplot)
```


```
#selection coefficients, with shaded central tendencies and 90% CIs
mcmc_areas(results, pars = c( paste0("b[",seq(1:3),"]"), paste0("q[",seq(1:3),"]"), paste0("qc[",seq(1:3),"]")))
```



340

341


```
#variance & corrs of RN parameters, mean and 50% CIs (dark line) and 90% CIs (light line)
mcmc_intervals(results, pars = c( paste0("V_P[",seq(1:3),"),"), "R[1,2]", "R[1,3]", "R[2,3]" ) )
```



Point estimates summarizing these posteriors can be quickly generated by summarizing the model.

```
#only first 17 parameters, round to ease interpretation
round(summary(results)$summary[1:16,], 2)
```

	##	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
344	## mu_0	0.00	0.00	0.04	-0.09	-0.03	0.00	0.03	0.09	890.59	1.00
345	## beta_x	-0.01	0.00	0.05	-0.11	-0.04	-0.01	0.03	0.09	900.31	1.00
346	## sigma_0	0.71	0.00	0.06	0.59	0.67	0.71	0.75	0.81	473.27	1.01
347	## W_0	0.97	0.01	0.11	0.76	0.90	0.96	1.04	1.20	222.47	1.03
348	## b[1]	-0.19	0.00	0.08	-0.36	-0.25	-0.19	-0.14	-0.04	343.58	1.01
349	## b[2]	-0.13	0.00	0.07	-0.28	-0.18	-0.13	-0.08	0.01	318.98	1.02
350	## b[3]	0.32	0.01	0.08	0.17	0.26	0.31	0.37	0.47	152.39	1.05
351	## q[1]	0.34	0.01	0.09	0.17	0.28	0.34	0.39	0.52	213.23	1.01
352	## q[2]	-0.23	0.01	0.08	-0.40	-0.28	-0.22	-0.17	-0.07	199.15	1.01
353	## q[3]	-0.21	0.01	0.08	-0.37	-0.26	-0.21	-0.16	-0.06	99.43	1.06
354	## qc[1]	-0.22	0.01	0.12	-0.45	-0.30	-0.22	-0.14	0.03	148.90	1.01
355	## qc[2]	0.24	0.01	0.13	-0.01	0.15	0.24	0.32	0.50	170.02	1.02
356	## qc[3]	-0.15	0.01	0.13	-0.42	-0.23	-0.15	-0.07	0.08	103.61	1.04
357	## sd_RN[1]	0.87	0.00	0.05	0.78	0.84	0.87	0.91	0.97	447.50	1.01
358	## sd_RN[2]	0.96	0.00	0.05	0.87	0.93	0.96	0.99	1.06	573.02	1.00
359	## sd_RN[3]	0.96	0.00	0.05	0.86	0.93	0.96	1.00	1.07	481.84	1.01
360											

The extracted posterior samples can also be manually plotted and summarized using base R functions. The `shinystan` package also provides a very helpful graphical user interface for looking at all aspects of model

363 fit and estimation. Running this code will open a new window in your internet browser for looking at the
364 model in greater detail.

```
library(shinystan)
launch_shinystan(results)
```

365 Hypothesis testing

366 MCMC not only facilitates sampling complex Bayesian models but also conducting straightforward and
367 direct forms of hypothesis testing. For example, if we want to know how much support there is for positive
368 linear and nonlinear selection effects p_+ , we simply need to calculate the proportion of the MCMC samples
369 for these parameters with positive magnitude, which approximates the area under the posterior distribution
370 providing support for positive effects.

```
#for each column, calculate probability of positive effect
apply(samples$b, 2, FUN = function(x) sum(x>0)/length(x) )
```

```
371 ## [1] 0.0075 0.0300 1.0000
```

```
apply(samples$q, 2, FUN = function(x) sum(x>0)/length(x) )
```

```
372 ## [1] 1.0000 0.0015 0.0045
```

```
apply(samples$qc, 2, FUN = function(x) sum(x>0)/length(x) )
```

```
373 ## [1] 0.040 0.967 0.109
```

374 Strong support for positive selection effects will approach 1, while strong support for negative selection effects
375 will approach 0. If the model is uncertain about the direction, the value will instead move toward 0.5. A
376 more general way to write the function is to test whatever direction the median effect size is in (so p_+ or p_-
377 depending on the location of the posterior).

```
#for each column, calculate probability of positive effect
apply(samples$b, 2, FUN = function(x) sum(sign(x) == sign(median(x)))/length(x) )
```

```
378 ## [1] 0.9925 0.9700 1.0000
```

```
apply(samples$q, 2, FUN = function(x) sum(sign(x) == sign(median(x)))/length(x) )
```

```
379 ## [1] 1.0000 0.9985 0.9955
```

```
apply(samples$qc, 2, FUN = function(x) sum(sign(x) == sign(median(x)))/length(x) )
```

```
380 ## [1] 0.960 0.967 0.891
```

381 We can also calculate quantile-based Bayesian credible intervals (CI). 90% CI are particularly useful for
382 interpretation. If the 90% CI excludes 0, then there is at least 0.95 probability in support of a directional
383 effect (i.e. ≤ 0.5 for an effect in the opposite direction).

```

#for each column, calculate quantile based CI
apply(samples$b, 2, FUN = function(x) quantile(x, c(0.05, 0.95)) ) #90% CI

384 ##
385 ##           [,1]           [,2]           [,3]
386 ##    5%  -0.33416102 -0.2575130  0.1912458
387 ##    95% -0.06632102 -0.0137624  0.4498332

apply(samples$q, 2, FUN = function(x) quantile(x, c(0.05, 0.95)) ) #90% CI

388 ##
389 ##           [,1]           [,2]           [,3]
390 ##    5%   0.1939239 -0.37042841 -0.34518754
391 ##    95%  0.4790374 -0.09467571 -0.08289907

apply(samples$qc, 2, FUN = function(x) quantile(x, c(0.05, 0.95)) ) #90% CI

392 ##
393 ##           [,1]           [,2]           [,3]
394 ##    5%  -0.40934611  0.02526734 -0.3676424
395 ##    95% -0.01150826  0.45137351  0.0496357

```

396 It is important to keep in mind that the size of CI interval is semi-arbitrary and so should not in itself be used
 397 as a discrete threshold for determining whether effects are ‘significant’ (McElreath (2020)). Although 0.95
 398 is a useful heuristic for designating clear evidence of an effect, discretizing this information into “significant”
 399 or “non-significant” is generally a waste of information. Put another way, these Bayesian hypothesis tests
 400 provides a continuous measure of evidence that should also be interpreted continuously. Much as the differ-
 401 ence between a significant and non-significant result is itself often not statistically significant (see McShane
 402 et al. 2019 for discussion), so too is the difference between e.g. a posterior probability of 0.93 and 0.97 not
 403 necessarily indicative of crossing a biologically or mathematically meaningful threshold. Thus, one should
 404 eschew the notion that a posterior probability <0.95 indicates “no evidence of an effect”, and instead get
 405 comfortable describing varying degrees of support (weak, moderate, and strong) for or against hypothesized
 406 effects. Any probability greater than 0.50 provides some (perhaps extremely weak) support for an effect,
 407 but most researchers would be uncomfortable to confidently assert empirical claims without much greater
 408 empirical support in their favor, e.g. evidence approaching only a 1/20 chance of an effect in the opposite
 409 direction (i.e. a posterior probability of 0.95). Thus, one should qualify the strength of evidence and interpret
 410 accordingly, rather than simply dichotomizing effects into two bins. Encouraging this Bayesian attitude by
 411 taking a continuous approach to statistical inference encourages researchers to put greater emphasis on effect
 412 sizes and multiple other metrics that collectively increase or decrease the overall biological significance of an
 413 empirical finding (McShane et al. 2019).

414 A variety of other hypotheses could also be easily tested for any parameter in the model. For instance,
 415 we could ask whether there is support for directional selection on intercepts being greater than directional
 416 selection on slopes among individuals.

```

sum(samples$b[,1] > samples$b[,2])/length(samples$b[,1])

417 ## [1] 0.2915

```

418 We can also quantify further pieces of information about the difference of these effect sizes to inform our
 419 inferences

```

diff = samples$b[,1] - samples$b[,3] #posterior of the difference in coefs
median(diff) #central tendency of difference

420 ## [1] -0.5001592

mad(diff) #median absolute deviation (robust SD) of expected difference

421 ## [1] 0.1195088

quantile(diff, c(0.05,0.95)) #90% CI of expected difference

422 ##          5%          95%
423 ## -0.7108768 -0.3191017

424 If one is so inclined, more robust null hypothesis tests can also be conducted within a Bayesian framework
425 by specifying a range of biologically trivial effect sizes. For example, on a standardized scale, regression
426 coefficients  $-0.10 < r < 0.10$  are extremely small and usually negligible effects. We might, therefore,
427 think of these as “trivial hypothesis” tests rather than null hypothesis tests per se. Comparing these trivial
428 hypothesis tests with directional hypothesis tests can provide further pieces of information. Consider the
429 RN parameter correlations

R=samples$R[,,] #3d array, 1 dim = samples, 2 dim = rows, 3 dim = columns

#directional hypothesis tests
sum(R[,1,2]<0)/length(R[,1,2]) #cor(intercept, slope)

430 ## [1] 0.71

sum(R[,1,3]>0)/length(R[,1,3]) #cor(intercept, residual)

431 ## [1] 0.515

sum(R[,2,3]>0)/length(R[,2,3]) #cor(slope, residual)

432 ## [1] 0.0645

#trivial hypothesis tests
sum(-0.1< R[,1,2] & R[,1,2] <0.1)/length(R[,1,2])

433 ## [1] 0.8005

sum(-0.1< R[,1,3] & R[,1,3] <0.1)/length(R[,1,3])

434 ## [1] 0.7665

```

```
sum(-0.1 < R[,2,3] & R[,2,3] < 0.1)/length(R[,2,3])
```

```
## [1] 0.3375
```

The directional and trivial hypothesis tests work together to inform our understanding of the direction and magnitude of the estimated correlations. While strong support is provided for non-trivial correlations among intercepts and slopes and residuals, the correlation among intercepts and residuals is both highly uncertain in its direction and largely overlapping the trivial effect size range, providing little evidence for an association of biologically meaningful magnitude

Calculate selection gradients

Analytic approach (Gaussian case)

We now want to calculate selection gradients from the estimated model. For the simple Gaussian case considered here, this is straightforward to accomplish analytically. Since we didn't mean-scale fitness prior to the analysis, we need to divide all gradients by mean fitness. For Gaussian measures, one can of course also mean-scale fitness prior to analysis. However, the sample mean will often be confounded by measurement error (Dingemanse, Araya-Ajoy, and Westneat 2021), motivating use of the estimated fitness intercept W_0 . In cases where adjusted effects are included, it is important to consider whether the intercept is equal to average fitness or instead conditional of a specific factor or covariate value. The quadratic selection effects for the squared trait values (i.e. not correlational selection effects) then need to be multiplied by 2 (Stinchcombe et al. 2008). This step could be circumvented by including $0.5 * q$ in the Stan model formula above. However, we have not done so here in keeping with the general model structure presented in the main text, which applies irrespective of the trait distribution. It is important that these arithmetic operations are done across each posterior sample to ensure that posterior uncertainty is retained across stages of analysis (Stinchcombe, Simonsen, and Blows 2014).

```
b = samples$b[,1:3]
q = samples$q[,1:3]
qc = samples$qc[,1:3]
```

```
beta = apply(b, 2, function(x) x / samples$W_0 )
colnames(beta) = c("mu", "beta", "sigma")
data.frame(median = apply(beta, 2, median),
           robust.SD = apply(beta, 2, mad),
           pd = apply(beta, 2, function(x)
             sum(sign(x) == sign(median(x))) / length(x)))
```

```
##           median robust.SD      pd
## mu      -0.1964688 0.08650714 0.9925
## beta    -0.1352696 0.07923651 0.9700
## sigma    0.3258650 0.08071981 1.0000
```

```
#doubling values (unless explicitly putting 0.5 * q in likelihood)
quad_diag = apply(q, 2, function(x) 2 * x / samples$W_0 )
colnames(quad_diag) = c("mu.mu", "beta.beta", "sigma.sigma")
data.frame(median = apply(quad_diag, 2, median),
           robust.SD = apply(quad_diag, 2, mad),
           pd = apply(quad_diag, 2, function(x)
             sum(sign(x) == sign(median(x))) / length(x)))
```

```

460 ##                median robust.SD      pd
461 ## mu.mu          0.7014768 0.2142443 1.0000
462 ## beta.beta      -0.4634167 0.1499078 0.9985
463 ## sigma.sigma    -0.4320118 0.1460994 0.9955

```

```

#correlational selection
quad_corr = apply(qc, 2, function(x) x / samples$W_0 )
colnames(quad_corr) = c("mu.beta", "mu.sigma", "beta.sigma")
data.frame(median = apply(quad_corr, 2, median),
           robust.SD = apply(quad_corr, 2, mad),
           pd = apply(quad_corr, 2, function(x)
             sum(sign(x) == sign(median(x))) / length(x)))

```

```

464 ##                median robust.SD      pd
465 ## mu.beta        -0.2264383 0.1256402 0.960
466 ## mu.sigma       0.2459766 0.1311841 0.967
467 ## beta.sigma     -0.1572348 0.1243206 0.891

```

468 Standardized gradients can also be calculated with the associated RN parameter variances and standard
 469 deviations (or population means).

```

#function for correlational gradients
pwSD = function(M) combn(ncol(M), 2, \(ij) as.matrix(M)[, ij[1]] * as.matrix(M)[, ij[2]])

sd.beta = beta * samples$sd_RN
sd.quad_diag = quad_diag * samples$sd_RN^2
sd.quad_corr = quad_corr * pwSD(samples$sd_RN)

```

470 The same techniques for summarizing and visualizing shown above can be used to plot and report these final
 471 results.

472 Numeric approach (general case)

473 In many cases, fitness residuals will not be well described by a Gaussian distribution, warranting an alter-
 474 native numeric approach. We demonstrate this here with our Gaussian dataset to show the compatibility
 475 of these approaches. See (Martin et al. 2025) for a more in-depth worked example of a complex model
 476 combining Gaussian, binomial, and cumulative-logit (ordinal) fitness models.

477 The first step codes the fitness function \bar{W} in R with the estimated parameters so that we can return
 478 the expected mean fitness (θ). We want the function to capture all posterior uncertainty, and so it is run
 479 over each iteration i of MCMC samples. The RN parameter values are denoted by a vector \mathbf{p} with as many
 480 elements as there are parameters under selection (so in this case, 3). Adapting this code to non-Gaussian
 481 distributions simply requires using the right link functions, e.g. `theta = logistic(...)` for a Binomial
 482 model with logit link or `theta = exp(...)` for a Poisson model with log link. The structure of the function
 483 should exactly match the likelihood function in the Stan model.

484 We then use the `numDeriv` package to calculate the selection gradients by simply taking the first and second
 485 partial derivatives of the function with respect to expected RN parameters across the population, which
 486 are by definition 0 due to mean centering. Lists are used to collect the results for each iteration before
 487 reorganizing.

```

library(numDeriv)

beta_num = list()
gamma_num = list()

for(i in 1:length(samples$W_0)){ #total posterior samples
  #Wbar = returns mean value of fitness function based on model likelihood
  #p[1]-[3] are placeholders for reaction norm intercepts, slopes, and residuals
  Wbar = function(p,i) {
    theta = (
      samples$W_0[i] + samples$W_0j[i, stan.dl$ind_z] +

      b[i,1] * p[1] +
      b[i,2] * p[2] +
      b[i,3] * p[3] +

      q[i,1] * p[1]^2 +
      q[i,2] * p[2]^2 +
      q[i,3] * p[3]^2 +

      qc[i,1] * p[1] * p[2] +
      qc[i,2] * p[1] * p[3] +
      qc[i,3] * p[2] * p[3]
    )

    return(mean(theta)) }

  p = rep(0,3)
  first.derivatives = grad(func = Wbar, x = p, i = i)
  second.derivatives = hessian(func = Wbar, x = p, i = i)
  denom = Wbar(p = p, i = i)
  beta_num[[i]] = first.derivatives/denom
  gamma_num[[i]] = second.derivatives/denom
}

#organize results
pars = c("mu", "beta", "sigma")
beta2 = do.call(rbind.data.frame, beta_num)
colnames(beta2) = pars

gamma_num2 = do.call(rbind.data.frame, lapply(gamma_num, reshape2::melt) )
gamma_num2 = gamma_num2[gamma_num2$Var1 <= gamma_num2$Var2,]
gamma_num2$Var1 = factor(pars[gamma_num2$Var1], levels = pars)
gamma_num2$Var2 = factor(pars[gamma_num2$Var2], levels = pars)
gamma_num2$vars = factor(paste0(gamma_num2$Var1,".", gamma_num2$Var2),
  levels = c("mu.mu", "beta.beta", "sigma.sigma",
    "mu.beta", "mu.sigma", "beta.sigma"))
quad_diag2 = droplevels(gamma_num2[gamma_num2$Var1==gamma_num2$Var2, ])
quad_diag2 = as.data.frame(split(quad_diag2$value, quad_diag2$vars))
quad_corr2 = droplevels(gamma_num2[gamma_num2$Var1!=gamma_num2$Var2, ])
quad_corr2 = as.data.frame(split(quad_corr2$value, quad_corr2$vars))

#summarize
apply(beta2, 2, median)

```

```

488 ##      mu      beta      sigma
489 ## -0.1956786 -0.1347033  0.3262571

```

```

apply(quad_diag2, 2, median)

```

```

490 ##      mu.mu      beta.beta sigma.sigma
491 ##   0.7019211   -0.4624542  -0.4317222

```

```

apply(quad_corr2, 2, median)

```

```

492 ##      mu.beta      mu.sigma beta.sigma
493 ## -0.2277783   0.2463497  -0.1574348

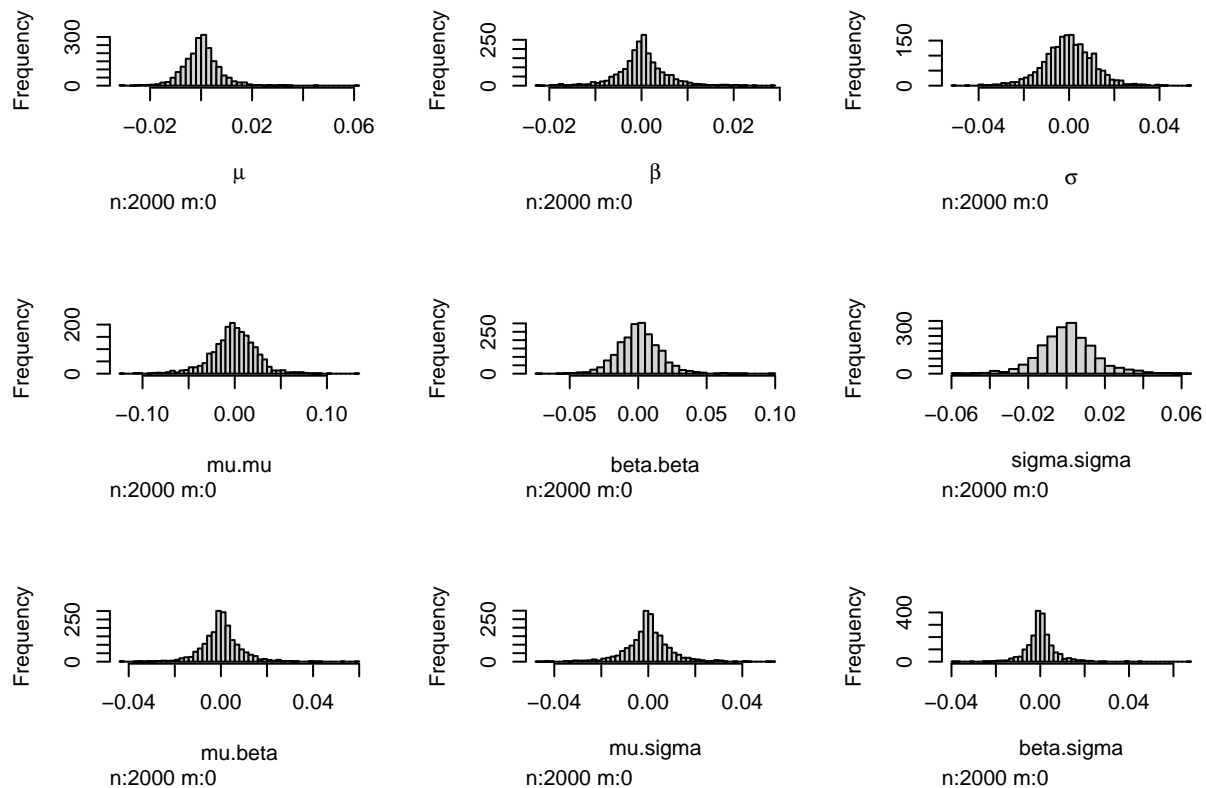
```

494 Very small, stochastic differences are expected between the analytic and numeric approaches, but the ex-
 495 pected difference is centered on 0, indicating convergence between methods.

```

library(Hmisc)
#compare
hist(cbind(beta2 - beta, quad_diag2 - quad_diag, quad_corr2 - quad_corr))

```



496
 497

References

- Carpenter, B., A. Gelman, M. D. Hoffman, D. Lee, B. Goodrich, M. Betancourt, and... A. Riddell. 2017. "Stan: A Probabilistic Programming Language." *Journal of Statistical Software* 74. <https://www.jstatsoft.org/article/view/v076i01>.
- Dingemanse, N. J., Y. G. Araya-Ajoy, and D. F. Westneat. 2021. "Most Published Selection Gradients Are Underestimated: Why This Is and How to Fix It." *Evolution* Early View.
- Hoffman, M. D., and A. Gelman. 2014. "The No-u-Turn Sampler: Adaptively Setting Path Lengths in Hamiltonian Monte Carlo." *Journal of Machine Learning Research* 15: 1593–623.
- Lemoine, N. P. 2019. "Moving Beyond Noninformative Priors: Why and How to Choose Weakly Informative Priors in Bayesian Analyses." *Oikos* 128. <https://onlinelibrary.wiley.com/doi/full/10.1111/oik.05985>.
- Link, W. A., and M. J. Eaton. 2012. "On Thinning of Chains in MCMC." *Methods in Ecology and Evolution* 3: 112–15.
- Martin, J. S., D. F. Westneat, S. Nakagawa, J. Schroeder, and T. Burke. 2025. "Measuring Selection on Reaction Norms: Lack's Principle and Plasticity in Clutch Size." *Evolution* qpap168.
- McElreath, R. 2020. *Statistical Rethinking: A Bayesian Course with Examples in r and Stan*. 2nd ed. CRC Press. <https://xcelab.net/rm/statistical-rethinking/>.
- McShane, B. B., D. Gal, A. Gelman, C. Robert, and J. L. Tackett. 2019. "Abandon Statistical Significance." *The American Naturalist* 73: 235–45.
- R Core Team. 2020. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org>.
- Stinchcombe, J. R., A. F. Agrawal, P. A. Hohenlohe, S. J. Arnold, and M. W. Blows. 2008. "Estimating Nonlinear Selection Gradients Using Quadratic Regression Coefficients: Double or Nothing?" *Evolution* 68. <https://onlinelibrary.wiley.com/doi/full/10.1111/evo.12321>.
- Stinchcombe, J. R., A. K. Simonsen, and M. W. Blows. 2014. "Estimating Uncertainty in Multivariate Responses to Selection." *Evolution* 68. <https://onlinelibrary.wiley.com/doi/full/10.1111/evo.12321>.